

Table of Contents

## CONV2D

---

**CLASS** `torch.nn.Conv2d(in_channels: int, out_channels: int, kernel_size: Union[T, Tuple[T, T]], stride: Union[T, Tuple[T, T]] = 1, padding: Union[T, Tuple[T, T]] = 0, dilation: Union[T, Tuple[T, T]] = 1, groups: int = 1, bias: bool = True, padding_mode: str = 'zeros')` [SOURCE]

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size  $(N, C_{in}, H, W)$  and output  $(N, C_{out}, H_{out}, W_{out})$  can be precisely described as:

$$\text{out}(N_i, C_{out_j}) = \text{bias}(C_{out_j}) + \sum_{k=0}^{C_{in}-1} \text{weight}(C_{out_j}, k) \star \text{input}(N_i, k)$$

where  $\star$  is the valid 2D [cross-correlation](#) operator,  $N$  is a batch size,  $C$  denotes a number of channels,  $H$  is a height of input planes in pixels, and  $W$  is width in pixels.

This module supports [TensorFloat32](#).

- `stride` controls the stride for the cross-correlation, a single number or a tuple.
- `padding` controls the amount of implicit zero-paddings on both sides for `padding` number of points for each dimension.
- `dilation` controls the spacing between the kernel points; also known as the *à trous* algorithm. It is harder to describe, but this [link](#) has a nice visualization of what `dilation` does.
- `groups` controls the connections between inputs and outputs. `in_channels` and `out_channels` must both be divisible by `groups`. For example,

- At `groups=1`, all inputs are convolved to all outputs.
- At `groups=2`, the operation becomes equivalent to having two conv layers side by side, each seeing half the input channels, and producing half the output channels, and both subsequently concatenated.
- At `groups= in_channels`, each input channel is convolved with its own set of filters, of size:  $\left\lfloor \frac{\text{out\_channels}}{\text{in\_channels}} \right\rfloor$ .

The parameters `kernel_size`, `stride`, `padding`, `dilation` can either be:

- a single `int` – in which case the same value is used for the height and width dimension
- a `tuple` of two `ints` – in which case, the first `int` is used for the height dimension, and the second `int` for the width dimension

### • NOTE

Depending of the size of your kernel, several (of the last) columns of the input might be lost, because it is a valid [cross-correlation](#), and not a full [cross-correlation](#). It is up to the user to add proper padding.

### • NOTE

When `groups==in_channels` and `out_channels==K*in_channels`, where  $K$  is a positive integer, this operation is also termed in literature as depthwise convolution.

In other words, for an input of size  $(N, C_{in}, H_{in}, W_{in})$ , a depthwise convolution with a depthwise multiplier  $K$ , can be constructed by arguments  $(in\_channels = C_{in}, out\_channels = C_{in} \times K, ..., groups = C_{in})$ .

### • NOTE

In some circumstances when using the CUDA backend with CuDNN, this operator may select a nondeterministic algorithm to increase performance. If this is undesirable, you can try to make the operation deterministic (potentially at a performance cost) by setting `torch.backends.cudnn.deterministic = True`. Please see the notes on [Reproducibility](#) for background.

---

### Parameters

- `in_channels`** (*int*) – Number of channels in the input image
- `out_channels`** (*int*) – Number of channels produced by the convolution
- `kernel_size`** (*int or tuple*) – Size of the convolving kernel
- `stride`** (*int or tuple, optional*) – Stride of the convolution. Default: 1

- **padding** (*int or tuple, optional*) – Zero-padding added to both sides of the input. Default: 0
- **padding\_mode** (*string, optional*) – 'zeros', 'reflect', 'replicate' or 'circular'. Default: 'zeros'
- **dilation** (*int or tuple, optional*) – Spacing between kernel elements. Default: 1
- **groups** (*int, optional*) – Number of blocked connections from input channels to output channels. Default: 1
- **bias** (*bool, optional*) – If `True`, adds a learnable bias to the output. Default: `True`

---

Shape:

- Input:  $(N, C_{in}, H_{in}, W_{in})$
- Output:  $(N, C_{out}, H_{out}, W_{out})$  where

$$H_{out} = \left\lfloor \frac{H_{in} + 2 \times \text{padding}[0] - \text{dilation}[0] \times (\text{kernel\_size}[0] - 1) - 1}{\text{stride}[0]} + 1 \right\rfloor$$

$$W_{out} = \left\lfloor \frac{W_{in} + 2 \times \text{padding}[1] - \text{dilation}[1] \times (\text{kernel\_size}[1] - 1) - 1}{\text{stride}[1]} + 1 \right\rfloor$$


---

Variables

- **~Conv2d.weight** (*Tensor*) – the learnable weights of the module of shape  $(\text{out\_channels}, \frac{\text{in\_channels}}{\text{groups}}, \text{kernel\_size}[0], \text{kernel\_size}[1])$ . The values of these weights are sampled from  $\mathcal{U}(-\sqrt{k}, \sqrt{k})$  where  $k = \frac{\text{groups}}{C_{in} * \prod_{i=0}^1 \text{kernel\_size}[i]}$
- **~Conv2d.bias** (*Tensor*) – the learnable bias of the module of shape  $(\text{out\_channels})$ . If `bias` is `True`, then the values of these weights are sampled from  $\mathcal{U}(-\sqrt{k}, \sqrt{k})$  where  $k = \frac{\text{groups}}{C_{in} * \prod_{i=0}^1 \text{kernel\_size}[i]}$

Examples

```
>>> # With square kernels and equal stride
>>> m = nn.Conv2d(16, 33, 3, stride=2)
>>> # non-square kernels and unequal stride and with padding
>>> m = nn.Conv2d(16, 33, (3, 5), stride=(2, 1), padding=(4, 2))
>>> # non-square kernels and unequal stride and with padding and dilation
>>> m = nn.Conv2d(16, 33, (3, 5), stride=(2, 1), padding=(4, 2), dilation=(3, 1))
>>> input = torch.randn(20, 16, 50, 100)
>>> output = m(input)
```

[< Previous](#)

[Next >](#)

© Copyright 2019, Torch Contributors.

Built with [Sphinx](#) using a [theme](#) provided by [Read the Docs](#).